

# Project

Isaac Attuah

Saturday, February 19th 2022

## Preparing Notebook

```
# Clear workspace
rm(list = ls())

# Set Working Directory
# setwd("./assignment_one")
# getwd()

# Load Libraries
library(ISLR2)
library(ggcorrplot)
```

```
## Loading required package: ggplot2
```

```
library(tidyverse)
```

```
## -- Attaching packages ----- tidyverse 1.3.1 --
```

```
## v tibble  3.1.6      v dplyr   1.0.7
## v tidyr   1.1.4      v stringr 1.4.0
## v readr   2.0.1      v forcats 0.5.1
## v purrr   0.3.4
```

```
## -- Conflicts ----- tidyverse_conflicts() --
```

```
## x dplyr::filter() masks stats::filter()
## x dplyr::lag()    masks stats::lag()
```

```
library(MASS)
```

```
##
```

```
## Attaching package: 'MASS'
```

```
## The following object is masked from 'package:dplyr':
```

```
##
```

```
##      select
```

```
## The following object is masked from 'package:ISLR2':  
##  
## Boston
```

```
library(boot)  
library(tree)
```

```
## Registered S3 method overwritten by 'tree':  
## method from  
## print.tree cli
```

```
library(randomForest)
```

```
## randomForest 4.7-1
```

```
## Type rfNews() to see new features/changes/bug fixes.
```

```
##  
## Attaching package: 'randomForest'
```

```
## The following object is masked from 'package:dplyr':  
##  
## combine
```

```
## The following object is masked from 'package:ggplot2':  
##  
## margin
```

```
# Load Dataset
```

```
onset_data <- read.csv("onset.csv")  
extra_data <- read.csv("armed.csv")
```

```
# Merge Datasets
```

```
war_data = merge(onset_data, extra_data, by = c("year", "gwno_a"))
```

```
# shuffle the dataframe by rows
```

```
# war_data <- war_data[sample(1:nrow(war_data)), ]
```

```
# Turn multiple columns to factor
```

```
cols <- c("gwno_a", "newconf", "onset1", "onset2", "onset3", "onset5", "onset10", "onset20", "intensity")  
war_data[cols] <- lapply(war_data[cols], as.factor)
```

```
# Clean and Attach Data
```

```
# war_data <- war_data[!names(war_data) %in% c("year_prev")]
```

```
war_data <- na.omit(war_data)
```

```
#attach(war_data)
```

```
# wd_merge <- merge(x=war_data, y=extra, by="gwno_a")
```

```
# war_data <- wd_merge
```

```
# View Data
```

```
head(war_data, 10)
```

```

##      year gwno_a abc      name newconf onset1 onset2 onset3 onset5 onset10
## 1  1946   145 BOL      Bolivia      1      1      1      1      1      1
## 2  1946   200 UKG United Kingdom      1      1      1      1      1      1
## 3  1946   210 NTH    Netherlands      1      1      1      1      1      1
## 4  1946   220 FRN      France      1      1      1      1      1      1
## 5  1946   220 FRN      France      1      1      1      1      1      1
## 6  1946   220 FRN      France      1      1      1      1      1      1
## 7  1946   220 FRN      France      1      1      1      1      1      1
## 8  1946   220 FRN      France      1      1      1      1      1      1
## 9  1946   220 FRN      France      1      1      1      1      1      1
## 10 1946   220 FRN      France      1      1      1      1      1      1
##      onset20 year_prev duration incompatibility intensity_level
## 1          1      1815      7              1              1
## 2          1      1815      7              0              0
## 3          1      1815      7              0              0
## 4          1      1815      7              0              0
## 5          1      1815      7              0              0
## 6          1      1815      7              0              0
## 7          1      1815      7              0              1
## 8          1      1815      7              0              0
## 9          1      1815      7              0              0
## 10         1      1815      7              0              0
##      cumulative_intensity ep_end
## 1              1      1
## 2              0      1
## 3              0      0
## 4              0      1
## 5              0      0
## 6              0      0
## 7              1      0
## 8              0      1
## 9              0      0
## 10             0      0

```

## Dataset Information

### General Information

```
message('Dimensions of Dataset')
```

```
## Dimensions of Dataset
```

```
dim(war_data)
```

```
## [1] 1040  17
```

```
message("Number of Rows ", nrow(war_data))
```

```
## Number of Rows 1040
```

```
message("Number of Columns ", ncol(war_data))
```

```
## Number of Columns 17
```

## Data Summary

```
summary(war_data)
```


```
##      year      gwno_a      abc      name      newconf
## Min.   :1946    750    :177 Length:1040 Length:1040    0:529
## 1st Qu.:1972    775    :165 Class :character Class :character 1:511
## Median :1991    530    : 77 Mode  :character Mode  :character
## Mean   :1987    220    : 46
## 3rd Qu.:2001    630    : 34
## Max.   :2017    365    : 30
##      (Other):511
## onset1 onset2 onset3 onset5 onset10 onset20 year_prev duration
## 0:  0  0:174  0:262  0:367  0:417  0:497 Min.   :1815 Min.   :1.00
## 1:1040 1:866 1:778 1:673 1:623 1:543 1st Qu.:1815 1st Qu.:2.00
##      Median :1950 Median :6.00
##      Mean   :1903 Mean   :4.84
##      3rd Qu.:1992 3rd Qu.:7.00
##      Max.   :2015 Max.   :7.00
##
## incompatibility intensity_level cumulative_intensity ep_end
## Min.   :0.0000  0:838          0:475          0:669
## 1st Qu.:0.0000  1:202          1:565          1:371
## Median :0.0000
## Mean   :0.3404
## 3rd Qu.:1.0000
## Max.   :3.0000
##
```

## Correlation

Since the values are categorical, we will resort to using

```
war_data_cor <- war_data[ , !names(war_data) %in% c("abc", "name", "year", "gwno_a", "onset1")]

library(ggcorrplot)
model.matrix(~0+., data=war_data_cor) %>%
  cor(use="pairwise.complete.obs") %>%
  ggcorrplot(show.diag = F, type="lower", lab=TRUE, lab_size=2)
```



## Dividing Into Training and Testing

```
# Divide data into training and testing (# 3)
# Examples for training = (0.80 * 683) = 546 entries
# Examples for testing test = (0.20 * 683) = 137 entries

set.seed(222)
sample_size = round(nrow(war_data)*.80) # setting sample size is 80%
index <- sample(seq_len(nrow(war_data)), size = sample_size)

train_better <- war_data[index, ]
test_better <- war_data[-index, ]

message("Number of Training Examples: ", nrow(train_better))

## Number of Training Examples: 832

message("Number of Testing Examples: ", nrow(test_better))

## Number of Testing Examples: 208
```

```

# train_better
# test_better

train_valid <- train_better[ , !names(train_better) %in% c("abc","name", "year", "gwno_a", "onset1", "y
test_valid <- test_better[ , !names(test_better) %in% c("abc","name", "year", "gwno_a", "onset1", "year
war_data_valid <- war_data[ , !names(test_better) %in% c("abc","name", "year", "gwno_a", "onset1")]

```

## Phase 1: Predicting War Outcome In 20 Years

### Model Testing

#### Logistic Regression Model

```

# Making model with all input variables

#There is not enough variation in onset1 so we will not include in the regression
glm.fits = glm(onset20 ~ newconf+onset2+onset3+onset5+onset10+duration+year_prev+duration+incompatibili
              data = train_better, family = binomial)

```

```
## Warning: glm.fit: algorithm did not converge
```

```

# glm.fits = glm(duration ~ onset2+onset3+onset5+onset10+onset20, data = train_better)

summary(glm.fits)

```

```

##
## Call:
## glm(formula = onset20 ~ newconf + onset2 + onset3 + onset5 +
##      onset10 + duration + year_prev + duration + incompatibility +
##      intensity_level + cumulative_intensity + ep_end, family = binomial,
##      data = train_better)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -2.409e-06 -2.409e-06  2.409e-06  2.409e-06  2.409e-06
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)   -7.970e+01  2.427e+06  0.000    1.000
## newconf1      -5.313e+01  2.257e+05  0.000    1.000
## onset21       -5.313e+01  9.401e+04 -0.001    1.000
## onset31       -5.313e+01  9.742e+04 -0.001    1.000
## onset51       -5.313e+01  1.040e+05 -0.001    1.000
## onset101      -5.313e+01  1.221e+05  0.000    1.000
## duration       5.313e+01  7.757e+04  0.001    0.999
## year_prev      2.635e-12  1.213e+03  0.000    1.000
## incompatibility 1.213e-10  2.415e+04  0.000    1.000

```

```
## intensity_level1      3.972e-10  3.893e+04  0.000  1.000
## cumulative_intensity1 1.121e-10  3.100e+04  0.000  1.000
## ep_end1              3.333e-12  2.654e+04  0.000  1.000
##
## (Dispersion parameter for binomial family taken to be 1)
##
## Null deviance: 1.1515e+03 on 831 degrees of freedom
## Residual deviance: 4.8269e-09 on 820 degrees of freedom
## AIC: 24
##
## Number of Fisher Scoring iterations: 25
```

```
# Make predictions based on model
glm.probs = predict(glm.fits, test_better, type="response")

# Initialize vector with 109 elements
glm.pred = rep(0, nrow(test_better))
# Assign 1 to probabilities > 0.5
glm.pred[glm.probs > .5] = 1

message('0 for no conflict, 1 for new conflict')
```

```
## 0 for no conflict, 1 for new conflict
```

```
message('Confusion Matrix')
```

```
## Confusion Matrix
```

```
# Confusion Matrix
table(glm.pred, test_better$onset20)
```

```
##
## glm.pred    0    1
##           0 101    0
##           1   0 107
```

```
# Test Error
message('Test Error Rate')
```

```
## Test Error Rate
```

```
mean(glm.pred != test_better$onset20)
```

```
## [1] 0
```

**LDA Model**

```
# Making model with all input variables
```

```
lda.fit=lda(onset20 ~ newconf+onset2+onset3+onset5+onset10+duration+year_prev+duration+incompatibility+  
            data = train_better)
```

```
## Warning in lda.default(x, grouping, ...): variables are collinear
```

```
lda.fit
```

```
## Call:
```

```
## lda(onset20 ~ newconf + onset2 + onset3 + onset5 + onset10 +  
##     duration + year_prev + duration + incompatibility + intensity_level +  
##     cumulative_intensity + ep_end, data = train_better)
```

```
##
```

```
## Prior probabilities of groups:
```

```
##      0      1  
## 0.4759615 0.5240385
```

```
##
```

```
## Group means:
```

```
##   newconf1  onset21  onset31  onset51  onset101  duration  year_prev  
## 0 0.0000000 0.6464646 0.479798 0.270202 0.1717172 2.568182 1988.338  
## 1 0.9288991 1.0000000 1.000000 1.000000 1.0000000 6.928899 1825.901  
## incompatibility intensity_level1 cumulative_intensity1 ep_end1  
## 0      0.3535354      0.1111111      0.6666667 0.3409091  
## 1      0.3394495      0.2454128      0.4197248 0.3784404
```

```
##
```

```
## Coefficients of linear discriminants:
```

```
##                               LD1  
## newconf1      4.3704336087  
## onset21      -0.1416472147  
## onset31      -0.1588930788  
## onset51      -0.1567057302  
## onset101      1.8061343812  
## duration      0.2094459456  
## year_prev      0.0001166536  
## incompatibility 0.0377394397  
## intensity_level1 -0.1753330702  
## cumulative_intensity1 0.1759232544  
## ep_end1      -0.0121873708
```

```
summary(lda.fit)
```

```
##      Length Class  Mode  
## prior      2    -none- numeric  
## counts      2    -none- numeric  
## means     22    -none- numeric  
## scaling    11    -none- numeric  
## lev         2    -none- character  
## svd          1    -none- numeric  
## N            1    -none- numeric  
## call         3    -none- call  
## terms        3    terms  call  
## xlevels      8    -none- list
```



```
lda.pred <- predict(lda.fit , test_better)
```

```
message('2 for benign, 4 for malignant')
```

```
## 2 for benign, 4 for malignant
```

```
message('Confusion Matrix')
```

```
## Confusion Matrix
```

```
# Confusion Matrix
```

```
lda.class <- lda.pred$class
```

```
table(lda.class, test_better$onset20)
```

```
##
```

```
## lda.class    0    1
```

```
##           0 101    1
```

```
##           1   0 106
```

```
message('Test Error Rate')
```



```
## Test Error Rate
```

```
# Test Error
```

```
mean(lda.class != test_better$onset20)
```

```
## [1] 0.004807692
```

```
plot(lda.fit)
```




## Linear Discriminants

## Decision Trees (Generic)

```
tree.onset20=tree(onset20 ~ newconf+onset2+onset3+onset5+onset10+duration+year_prev+duration+incompatib
summary(tree.onset20)
```

```
##
## Classification tree:
## tree(formula = onset20 ~ newconf + onset2 + onset3 + onset5 +
##       onset10 + duration + year_prev + duration + incompatibility +
##       intensity_level + cumulative_intensity + ep_end, data = war_data_valid)
## Variables actually used in tree construction:
## [1] "duration"
## Number of terminal nodes:  2
## Residual mean deviance:  0 = 0 / 1038
## Misclassification error rate: 0 = 0 / 1040
```

```
plot(tree.onset20)
text(tree.onset20, pretty = 0, cex=0.75)
```



## Decision Trees (With Training & Testing)

```

# Train using training set
tree.onset20=tree(onset20 ~ newconf+onset2+onset3+onset5+onset10+duration+incompatibility+intensity_level)

# Test on test set using predict()
# type="class" to return the class prediction
tree.pred=predict(tree.onset20,test_valid,type="class")

# Confusion matrix
conf.matrix <- table(tree.pred,test_valid$onset20)
conf.matrix

##
## tree.pred    0    1
##           0 101    0
##           1   0 107

# Accuracy on test set
(conf.matrix[1,1] + conf.matrix[2,2])/(conf.matrix[1,1] + conf.matrix[2,2] + conf.matrix[1,2] + conf.matrix[2,1])

## [1] 1

```


## Regression Trees

```
set.seed(1)

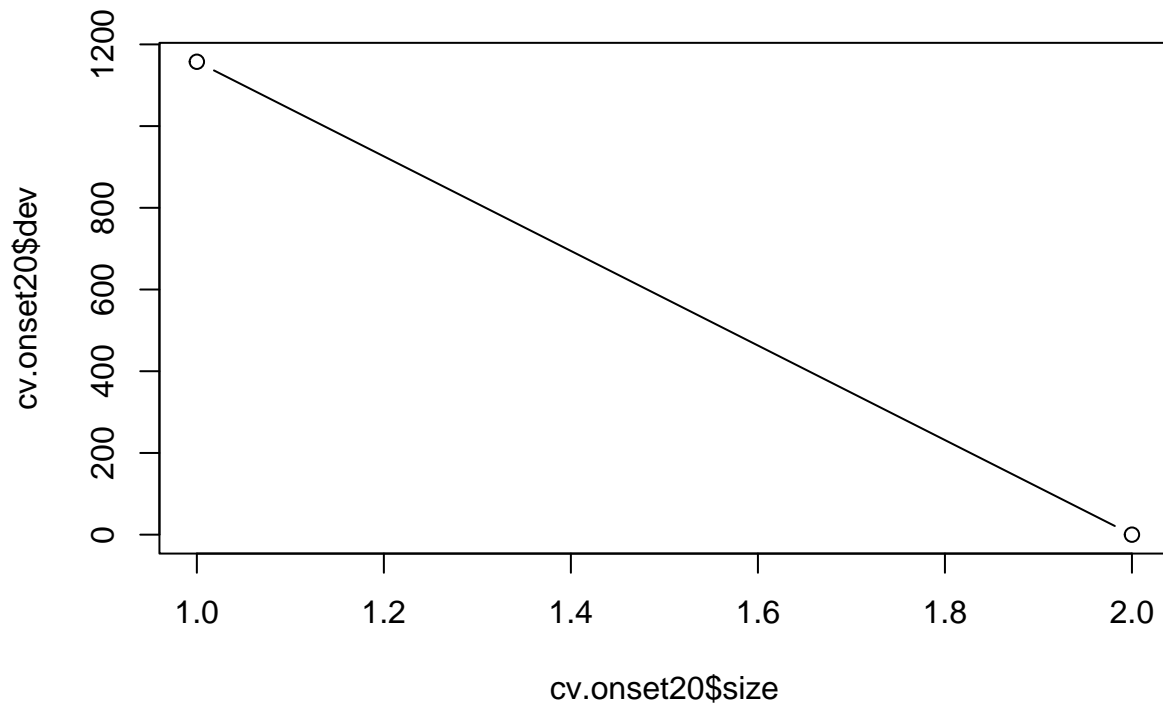
tree.onset20=tree(onset20 ~ newconf+onset2+onset3+onset5+onset10+duration+incompatibility+intensity_level,
# Only a few of the variables were used in constructing the tree
# lstat: percentage of individuals with lower socioeconomic status
summary(tree.onset20)

##
## Classification tree:
## tree(formula = onset20 ~ newconf + onset2 + onset3 + onset5 +
##       onset10 + duration + incompatibility + intensity_level +
##       cumulative_intensity + ep_end, data = train_valid)
## Variables actually used in tree construction:
## [1] "duration"
## Number of terminal nodes:  2
## Residual mean deviance:  0 = 0 / 830
## Misclassification error rate: 0 = 0 / 832

# Plot the tree
# Lower values of lstat correspond to more expensive houses
plot(tree.onset20)
text(tree.onset20,pretty=0,cex=0.75)
```




```
# cv.tree() to determine whether pruning improves performance
cv.onset20=cv.tree(tree.onset20)
# It doesn't seem to be the case
plot(cv.onset20$size,cv.onset20$dev,type="b")
```



```
# prune.tree(): function to prune to be used in case we wanted to prune the tree
prune.onset20=prune.tree(tree.onset20,best=5)
```

```
## Warning in prune.tree(tree.onset20, best = 5): best is bigger than tree size
```

```
plot(prune.onset20)
text(prune.onset20,pretty=0,cex=0.75)
```



```
# Predicting based on CV results (i.e., use the unpruned tree)
yhat=predict(tree.onset20,newdata=test_valid)
```

```
# plot(yhat,test_valid$onset20)
# abline(0,1)
# Test error
mse=mean((yhat-test_valid$onset20)^2)
```

```
## Warning in Ops.factor(yhat, test_valid$onset20): '-' not meaningful for factors
```

```
mse
```

```
## [1] NA
```

```
# This model leads to test predictions that are within around $5-6K of the true
# median home value for the suburb
sqrt(mse)
```

```
## [1] NA
```

## Random Forests

```

# By default randomForest() uses m=p/3 for regression and m=sqrt(p) for classification
# Let's try m=6
set.seed(1)
rf=randomForest(onset20 ~ newconf+onset2+onset3+onset5+onset10+duration+incompatibility+intensity_level,
yhat.rf = predict(tree.onset20,newdata=test_valid)

mean((yhat.rf-as.integer(test_valid$onset20))^2)

```

```
## [1] 1.528846
```

```

# importance(): view the importance of each variable
# %IncMSE: mean decrease of accuracy in predictions on the OOB samples when a
# given variable is excluded from the model
# IncNodeImpurity: total decrease in node impurity that results from splits over
# that variable, averaged over all trees (RSS in regr. vs. deviance in class.)
importance(rf)

```

```

##              0              1 MeanDecreaseAccuracy
## newconf      15.6892516  4.7108616      15.9759772
## onset2       0.0000000  2.0489485       2.0521178
## onset3       0.0000000  3.5784432       3.5878985
## onset5      -1.4169759  5.8180605       5.8190757
## onset10      1.5221103 10.3675345      10.3795729
## duration     39.6159791 32.5857121      44.0281649
## incompatibility 0.8375265 -1.3084633      -0.8899624
## intensity_level 0.3354407 -0.7267548      -0.4953894
## cumulative_intensity 1.8979866  2.6486145      3.1606838
## ep_end       1.7834743  0.5808985       1.5507581
##              MeanDecreaseGini
## newconf      115.44478476
## onset2       0.73073065
## onset3       3.43096536
## onset5      14.44642774
## onset10     46.09759296
## duration    233.43477076
## incompatibility 0.07150790
## intensity_level 0.04697760
## cumulative_intensity 0.16936079
## ep_end      0.04261412


```

```

# varImpPlot(): Variance importance plot
varImpPlot(rf)

```

rf



## Other Models

- Penalized Logistic Regression -plr
- Conditional Inference Random Forest -cforest
- Random Forest - rf
- Bayesian Generalized Linear Model -bayesglm
- Boosted Generalized Additive Model - gamboost
- Support Vector Machines with Linear Kernel - svmLinear

```
library(caret)
```

```
## Loading required package: lattice
```

```
##
```

```
## Attaching package: 'lattice'
```

```
## The following object is masked from 'package:boot':
```

```
##
```

```
## melanoma
```



```

##
## Attaching package: 'caret'

## The following object is masked from 'package:purrr':
##
## lift

#specify the cross-validation method
ctrl <- trainControl(method = "cv")

#fit a regression model and use LOOCV to evaluate performance
model <- train(onset20~newconf+onset2+onset3+onset5+onset10, data = train_better, method = "knn", trCon

#view summary of LOOCV
print(model)

## k-Nearest Neighbors
##
## 832 samples
## 5 predictor
## 2 classes: '0', '1'
##
## No pre-processing
## Resampling: Cross-Validated (10 fold)
## Summary of sample sizes: 749, 748, 748, 749, 749, 748, ...
## Resampling results across tuning parameters:
##
## k Accuracy Kappa
## 5 0.9627636 0.9256512
## 7 0.9627636 0.9256512
## 9 0.9627636 0.9256512
##
## Accuracy was used to select the optimal model using the largest value.
## The final value used for the model was k = 9.

predictions <- predict(model, test_better, type="raw")

message('0 for no conflict, 1 for new conflict')

## 0 for no conflict, 1 for new conflict

message('Confusion Matrix')

## Confusion Matrix

# Confusion Matrix
table(predictions,test_better$onset20)

##
## predictions 0 1
## 0 101 1
## 1 0 106

```

```
# Test Error
message('Test Error Rate')
```

```
## Test Error Rate
```

```
mean(predictions!=test_better$onset20)
```

```
## [1] 0.004807692
```

```
confusionMatrix(data = predict(model, test_better), test_better$onset20)
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction    0    1
##           0 101    1
##           1   0 106
##
##           Accuracy : 0.9952
##           95% CI : (0.9735, 0.9999)
##           No Information Rate : 0.5144
##           P-Value [Acc > NIR] : <2e-16
##
##           Kappa : 0.9904
##
##  Mcnemar's Test P-Value : 1
##
##           Sensitivity : 1.0000
##           Specificity : 0.9907
##           Pos Pred Value : 0.9902
##           Neg Pred Value : 1.0000
##           Prevalence : 0.4856
##           Detection Rate : 0.4856
##           Detection Prevalence : 0.4904
##           Balanced Accuracy : 0.9953
##
##           'Positive' Class : 0
##
```

## Phase 2: Predicting the End Of War

### Logistic Regression Model

```
# Making model with all input variables
```

```
#There is not enough variation in onset1 so we will not include in the regression
```

```
glm.fits = glm(ep_end ~ newconf+onset2+onset3+onset5+onset10+duration+year_prev+duration+incompatibilit;
              data = train_better, family = binomial)
```

```
# glm.fits = glm(duration ~ onset2+onset3+onset5+onset10+onset20, data = train_better)

summary(glm.fits)
```

```
##
## Call:
## glm(formula = ep_end ~ newconf + onset2 + onset3 + onset5 + onset10 +
##      duration + year_prev + duration + incompatibility + intensity_level +
##      cumulative_intensity + onset20, family = binomial, data = train_better)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -1.4624  -0.9715  -0.7358   1.2683   1.8536
##
## Coefficients: (1 not defined because of singularities)
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)   -4.276602   15.002399  -0.285   0.7756
## newconf1       0.712251    1.421590   0.501   0.6164
## onset21       0.298784    0.599713   0.498   0.6183
## onset31       0.315113    0.613570   0.514   0.6076
## onset51       0.247213    0.644233   0.384   0.7012
## onset101     -0.731827    0.766924  -0.954   0.3400
## duration     -0.055770    0.502538  -0.111   0.9116
## year_prev      0.001991    0.007498   0.266   0.7905
## incompatibility 0.310721    0.143585   2.164   0.0305 *
## intensity_level1 0.013462    0.246219   0.055   0.9564
## cumulative_intensity1 -0.957326    0.186151  -5.143 2.71e-07 ***
## onset201             NA          NA      NA      NA
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 1087.8  on 831  degrees of freedom
## Residual deviance: 1036.0  on 821  degrees of freedom
## AIC: 1058
##
## Number of Fisher Scoring iterations: 4
```

```
# Make predictions based on model
glm.probs = predict(glm.fits,test_better, type="response")
```

```
## Warning in predict.lm(object, newdata, se.fit, scale = 1, type = if (type == :
## prediction from a rank-deficient fit may be misleading
```

```
# Initialize vector with 109 elements
glm.pred = rep(0, nrow(test_better))
# Assign 1 to probabilities > 0.5
glm.pred[glm.probs > .5]=1

message('0 for no conflict, 1 for new conflict')
```

```
## 0 for no conflict, 1 for new conflict
```

```
message('Confusion Matrix')
```

```
## Confusion Matrix
```

```
# Confusion Matrix
```

```
table(glm.pred, test_better$ep_end)
```

```
##
```

```
## glm.pred    0    1
```

```
##           0 127  55
```

```
##           1  10  16
```

```
# Test Error
```

```
message('Test Error Rate')
```

```
## Test Error Rate
```

```
mean(glm.pred != test_better$ep_end)
```

```
## [1] 0.3125
```

## LDA Model

```
# Making model with all input variables
```

```
lda.fit = lda(ep_end ~ newconf + onset2 + onset3 + onset5 + onset10 + duration + year_prev + duration + incompatibility + intensity_level + cumulative_intensity, data = train_better)
```

```
## Warning in lda.default(x, grouping, ...): variables are collinear
```

```
lda.fit
```

```
## Call:
```

```
## lda(ep_end ~ newconf + onset2 + onset3 + onset5 + onset10 + duration +
```

```
##   year_prev + duration + incompatibility + intensity_level +
```

```
##   cumulative_intensity + onset20, data = train_better)
```

```
##
```

```
## Prior probabilities of groups:
```

```
##           0           1
```

```
## 0.6394231 0.3605769
```

```
##
```

```
## Group means:
```

```
##   newconf1  onset21  onset31  onset51  onset101  duration  year_prev
```

```
## 0 0.4661654 0.8120301 0.7293233 0.6390977 0.5996241 4.755639 1906.714
```

```
## 1 0.5233333 0.8666667 0.7933333 0.6766667 0.6166667 5.026667 1897.010
```

```
##   incompatibility  intensity_level1  cumulative_intensity1  onset201
```

```
## 0           0.3214286           0.2067669           0.6184211 0.5093985
```

```
## 1           0.3900000           0.1366667           0.3933333 0.5500000
```

```
##
```

```
## Coefficients of linear discriminants:
##                               LD1
## newconf1                     1.143387203
## onset21                      0.408209147
## onset31                      0.492529631
## onset51                      0.361322729
## onset101                    -1.517663352
## duration                     0.021684369
## year_prev                    0.003584622
## incompatibility              0.595685678
## intensity_level1             0.010460217
## cumulative_intensity1       -1.854945518
## onset201                    -0.107028235
```

```
summary(lda.fit)
```

```
##           Length Class  Mode
## prior      2      -none- numeric
## counts     2      -none- numeric
## means     22      -none- numeric
## scaling   11      -none- numeric
## lev        2      -none- character
## svd         1      -none- numeric
## N           1      -none- numeric
## call        3      -none- call
## terms       3      terms  call
## xlevels     8      -none- list
```

```
lda.pred <- predict(lda.fit , test_better)
```

```
message('2 for benign, 4 for malignant')
```

```
## 2 for benign, 4 for malignant
```

```
message('Confusion Matrix')
```

```
## Confusion Matrix
```

```
# Confusion Matrix
lda.class <- lda.pred$class
table(lda.class, test_better$onset20)
```

```
##
## lda.class  0  1
##           0 88 93
##           1 13 14
```



```
message('Test Error Rate')
```

```
## Test Error Rate
```

```
# Test Error
mean(lda.class != test_better$onset20)
```

```
## [1] 0.5096154
```

```
plot(lda.fit)
```



Linear Discriminants


Decision Trees (Generic)

```
tree.ep_end=tree(ep_end ~ newconf+onset2+onset3+onset5+onset10+duration+year_prev+duration+incompatibil.
summary(tree.ep_end)
```

```
##
## Classification tree:
## tree(formula = ep_end ~ newconf + onset2 + onset3 + onset5 +
##       onset10 + duration + year_prev + duration + incompatibility +
##       intensity_level + cumulative_intensity + ep_end + onset20,
##       data = war_data_valid)
## Variables actually used in tree construction:
```

```
## [1] "ep_end"
## Number of terminal nodes: 2
## Residual mean deviance: 0 = 0 / 1038
## Misclassification error rate: 0 = 0 / 1040
```

```
plot(tree.ep_end)
text(tree.onset20, pretty = 0, cex=0.75)
```



## Decision Trees (With Training & Testing)

```
# Train using training set
tree.ep_end=tree(ep_end ~ newconf+onset2+onset3+onset5+onset10+duration+incompatibility+intensity_level)

# Test on test set using predict()
# type="class" to return the class prediction
tree.ep_end=predict(tree.onset20,test_valid,type="class")

# Confusion matrix
conf.matrix <- table(tree.ep_end,test_valid$ep_end)
conf.matrix

##
## tree.ep_end  0  1
```

```
##          0 70 31
##          1 67 40
```

```
# Accuracy on test set
```

```
(conf.matrix[1,1] + conf.matrix[2,2])/(conf.matrix[1,1] + conf.matrix[2,2] + conf.matrix[1,2] + conf.matrix[2,1])
```

```
## [1] 0.5288462
```

## Regression Trees

```
set.seed(1)
```

```
tree.onset20=tree(ep_end ~ newconf+onset2+onset3+onset5+onset10+duration+incompatibility+intensity_level)
```

```
# Only a few of the variables were used in constructing the tree
```

```
# lstat: percentage of individuals with lower socioeconomic status
```

```
summary(tree.onset20)
```

```
##
```

```
## Classification tree:
```

```
## tree(formula = ep_end ~ newconf + onset2 + onset3 + onset5 +
```

```
##      onset10 + duration + incompatibility + intensity_level +
```

```
##      cumulative_intensity + onset20, data = train_valid)
```

```
## Variables actually used in tree construction:
```

```
## [1] "cumulative_intensity"
```

```
## Number of terminal nodes: 2
```

```
## Residual mean deviance: 1.263 = 1049 / 830
```

```
## Misclassification error rate: 0.3606 = 300 / 832
```


```
# Plot the tree
```

```
# Lower values of lstat correspond to more expensive houses
```


```
plot(tree.onset20)
```

```
text(tree.onset20,pretty=0,cex=0.75)
```






```
# cv.tree() to determine whether pruning improves performance  
cv.onset20=cv.tree(tree.onset20)  
# It doesn't seem to be the case  
plot(cv.onset20$size,cv.onset20$dev,type="b")
```



```
# prune.tree(): function to prune to be used in case we wanted to prune the tree  
prune.onset20=prune.tree(tree.onset20,best=5)
```

```
## Warning in prune.tree(tree.onset20, best = 5): best is bigger than tree size
```

```
plot(prune.onset20)  
text(prune.onset20,pretty=0,cex=0.75)
```



```
# Predicting based on CV results (i.e., use the unpruned tree)
yhat=predict(tree.onset20,newdata=test_valid)
```

```
# plot(yhat,test_valid$onset20)
# abline(0,1)
# Test error
mse=mean((yhat-test_valid$onset20)^2)
```

```
## Warning in Ops.factor(yhat, test_valid$onset20): '-' not meaningful for factors
```

```
mse
```

```
## [1] NA
```

```
# This model leads to test predictions that are within around $5-6K of the true
# median home value for the suburb
sqrt(mse)
```

```
## [1] NA
```

## Random Forests

```

# By default randomForest() uses m=p/3 for regression and m=sqrt(p) for classification
# Let's try m=6
set.seed(1)
rf=randomForest(ep_end ~ newconf+onset2+onset3+onset5+onset10+duration+incompatibility+intensity_level+
yhat.rf = predict(tree.onset20,newdata=test_valid)

mean((yhat.rf-as.integer(test_valid$onset20))^2)

```

```
## [1] 1.31077
```

```

# importance(): view the importance of each variable
# %IncMSE: mean decrease of accuracy in predictions on the OOB samples when a
# given variable is excluded from the model
# IncNodeImpurity: total decrease in node impurity that results from splits over
# that variable, averaged over all trees (RSS in regr. vs. deviance in class.)
importance(rf)

```

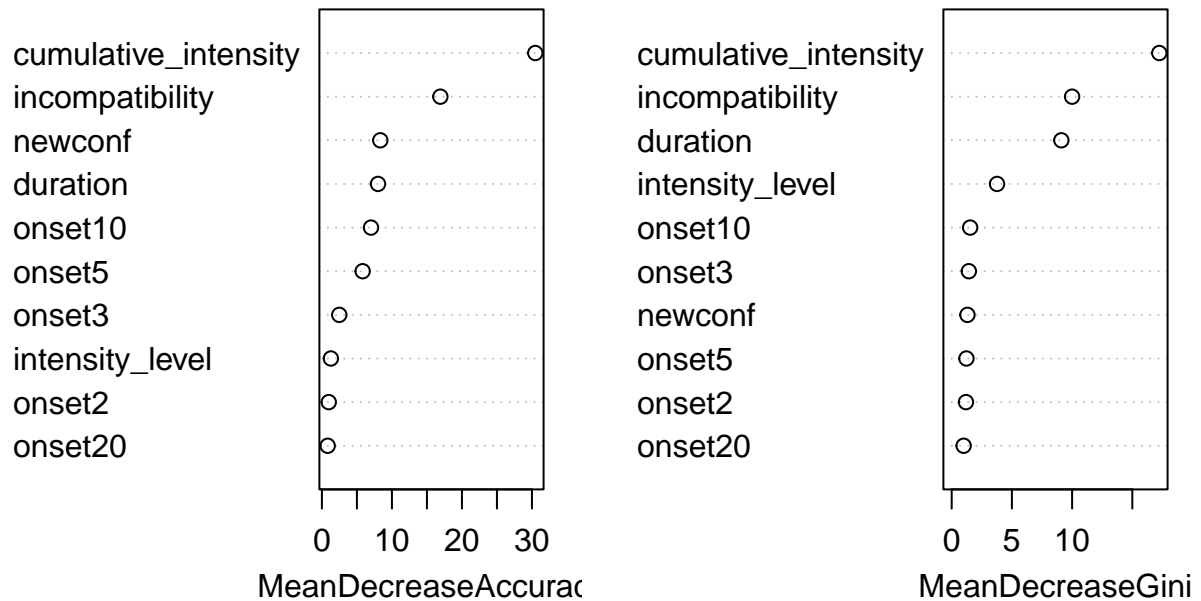
|                         | 0         | 1         | MeanDecreaseAccuracy | MeanDecreaseGini |
|-------------------------|-----------|-----------|----------------------|------------------|
| ## newconf              | 9.201656  | -6.007442 | 8.3412216            | 1.3093954        |
| ## onset2               | 1.955090  | -1.906697 | 0.9790256            | 1.1856673        |
| ## onset3               | 3.515001  | -2.485181 | 2.4813395            | 1.4251111        |
| ## onset5               | 9.347825  | -9.814094 | 5.8210619            | 1.2251029        |
| ## onset10              | 9.410704  | -7.425381 | 7.0121900            | 1.5338187        |
| ## duration             | 12.534343 | -9.563917 | 8.0113211            | 9.1080641        |
| ## incompatibility      | 19.804919 | 4.511146  | 16.9176142           | 9.9989336        |
| ## intensity_level      | -3.112003 | 5.660214  | 1.2834327            | 3.7662430        |
| ## cumulative_intensity | 18.746861 | 22.425895 | 30.4773404           | 17.2397718       |
| ## onset20              | 1.888401  | -1.411583 | 0.8251728            | 0.9932799        |

```

# varImpPlot(): Variance importance plot
varImpPlot(rf)

```

rf



## Other Models

- Penalized Logistic Regression -plr
- Conditional Inference Random Forest -cforest
- Random Forest - rf
- Bayesian Generalized Linear Model -bayesglm
- Boosted Generalized Additive Model - gamboost
- Support Vector Machines with Linear Kernel - svmLinear

```
library(caret)

#specify the cross-validation method
ctrl <- trainControl(method = "cv")

#fit a regression model and use LOOCV to evaluate performance
model <- train(ep_end ~ newconf+onset2+onset3+onset5+onset10+duration+incompatibility+intensity_level+c

#view summary of LOOCV
print(model)
```

## k-Nearest Neighbors

```
##
## 832 samples
## 10 predictor
## 2 classes: '0', '1'
##
## No pre-processing
## Resampling: Cross-Validated (10 fold)
## Summary of sample sizes: 749, 749, 748, 749, 748, 749, ...
## Resampling results across tuning parameters:
##
## k Accuracy Kappa
## 5 0.6453815 0.1190491
## 7 0.6478772 0.1246593
## 9 0.6538439 0.1273210
##
## Accuracy was used to select the optimal model using the largest value.
## The final value used for the model was k = 9.
```

```
predictions <- predict(model, test_better, type="raw")
message('0 for no conflict, 1 for new conflict')
```

```
## 0 for no conflict, 1 for new conflict
```

```
message('Confusion Matrix')
```

```
## Confusion Matrix
```

```
# Confusion Matrix
table(predictions, test_better$ep_end)
```

```
##
## predictions    0    1
##              0 130  55
##              1   7  16
```

```
# Test Error
message('Test Error Rate')
```

```
## Test Error Rate
```

```
mean(predictions != test_better$ep_end)
```

```
## [1] 0.2980769
```

```
confusionMatrix(data = predict(model, test_better), test_better$onset20)
```

```
## Confusion Matrix and Statistics
##
```

```
##           Reference
## Prediction  0   1
##           0 92 93
##           1   9 14
##
##           Accuracy : 0.5096
##           95% CI : (0.4396, 0.5794)
##           No Information Rate : 0.5144
##           P-Value [Acc > NIR] : 0.5826
##
##           Kappa : 0.0408
##
## Mcnemar's Test P-Value : <2e-16
##
##           Sensitivity : 0.9109
##           Specificity : 0.1308
##           Pos Pred Value : 0.4973
##           Neg Pred Value : 0.6087
##           Prevalence : 0.4856
##           Detection Rate : 0.4423
##           Detection Prevalence : 0.8894
##           Balanced Accuracy : 0.5209
##
##           'Positive' Class : 0
##
```

## Phase 3: Predicting War Susceptibility

### Logistic Regression Model

```
# Making model with all input variables

#There is not enough variation in onset1 so we will not include in the regression
glm.fits = glm(intensity_level ~ newconf+onset2+onset3+onset5+onset10+onset20+duration+year_prev+durati
              data = train_better, family = binomial)

# glm.fits = glm(duration ~ onset2+onset3+onset5+onset10+onset20, data = train_better)

summary(glm.fits)

##
## Call:
## glm(formula = intensity_level ~ newconf + onset2 + onset3 + onset5 +
##      onset10 + onset20 + duration + year_prev + duration + cumulative_intensity +
##      ep_end + incompatibility, family = binomial, data = train_better)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -1.69725  -0.42526  -0.00006  -0.00002   2.61536
##
## Coefficients: (1 not defined because of singularities)
```

```
##               Estimate Std. Error z value Pr(>|z|)
## (Intercept)    60.03491   836.96455   0.072 0.942817
## newconf1      -3.82102    1.84333  -2.073 0.038182 *
## onset21        0.77111    0.53679   1.437 0.150851
## onset31       -0.18438    0.54513  -0.338 0.735192
## onset51        0.69125    0.59928   1.153 0.248721
## onset101      -1.07498    0.61835  -1.738 0.082131 .
## onset201      -0.77004    0.74490  -1.034 0.301253
## duration              NA          NA      NA      NA
## year_prev       -0.04172    0.01152  -3.621 0.000293 ***
## cumulative_intensity1 20.60617  836.65014   0.025 0.980351
## ep_end1        -0.09638    0.27060  -0.356 0.721723
## incompatibility   0.24231    0.21655   1.119 0.263148
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##    Null deviance: 788.15  on 831  degrees of freedom
## Residual deviance: 434.79  on 821  degrees of freedom
## AIC: 456.79
##
## Number of Fisher Scoring iterations: 19
```

```
# Make predictions based on model
glm.probs = predict(glm.fits,test_better, type="response")
```

```
## Warning in predict.lm(object, newdata, se.fit, scale = 1, type = if (type == :
## prediction from a rank-deficient fit may be misleading
```

```
# Initialize vector with 109 elements
glm.pred = rep(0, nrow(test_better))
# Assign 1 to probabilities > 0.5
glm.pred[glm.probs > .5]=1

message('0 for no conflict, 1 for new conflict')
```

```
## 0 for no conflict, 1 for new conflict
```

```
message('Confusion Matrix')
```

```
## Confusion Matrix
```

```
# Confusion Matrix
table(glm.pred,test_better$intensity_level)
```

```
##
## glm.pred    0    1
##           0 142  14
##           1  15  37
```



```
# Test Error
message('Test Error Rate')
```

```
## Test Error Rate
```

```
mean(glm.pred!=test_better$intensity_level)
```

```
## [1] 0.1394231
```

## LDA Model

```
# Making model with all input variables
lda.fit=lda(intensity_level ~ newconf+onset2+onset3+onset5+onset10+onset20+duration+year_prev+duration+
            data = train_better)
```

```
## Warning in lda.default(x, grouping, ...): variables are collinear
```

```
lda.fit
```

```
## Call:
## lda(intensity_level ~ newconf + onset2 + onset3 + onset5 + onset10 +
##      onset20 + duration + year_prev + duration + cumulative_intensity +
##      ep_end + incompatibility, data = train_better)
##
## Prior probabilities of groups:
##      0      1
## 0.8185096 0.1814904
##
## Group means:
##      newconf1  onset21  onset31  onset51  onset101  onset201  duration
## 0 0.4419971 0.8061674 0.7224670 0.6138032 0.5697504 0.4831131 4.637298
## 1 0.6887417 0.9470199 0.8874172 0.8278146 0.7682119 0.7086093 5.827815
##      year_prev cumulative_intensity1  ep_end1 incompatibility
## 0 1911.604      0.4346549 0.3803231      0.3274596
## 1 1865.384      1.0000000 0.2715232      0.4304636
##
## Coefficients of linear discriminants:
##
##              LD1
## newconf1      -1.910244432
## onset21        0.382364337
## onset31        0.130624856
## onset51        0.546471805
## onset101      -0.509113280
## onset201      -0.530473192
## duration      -0.091087508
## year_prev     -0.023597594
## cumulative_intensity1 2.289064670
## ep_end1        0.003104976
## incompatibility 0.130908183
```

```
summary(lda.fit)
```

```
##           Length Class  Mode
## prior      2      -none- numeric
## counts     2      -none- numeric
## means     22      -none- numeric
## scaling   11      -none- numeric
## lev        2      -none- character
## svd         1      -none- numeric
## N           1      -none- numeric
## call        3      -none- call
## terms       3      terms  call
## xlevels     8      -none- list
```

```
lda.pred <- predict(lda.fit , test_better)
```

```
message('2 for benign, 4 for malignant')
```

```
## 2 for benign, 4 for malignant
```

```
message('Confusion Matrix')
```

```
## Confusion Matrix
```

```
# Confusion Matrix
```

```
lda.class <- lda.pred$class
table(lda.class, test_better$ep_end)
```

```
##
## lda.class  0  1
##           0 97 59
##           1 40 12
```

```
message('Test Error Rate')
```



```
## Test Error Rate
```

```
# Test Error
```

```
mean(lda.class != test_better$ep_end)
```

```
## [1] 0.4759615
```

```
plot(lda.fit)
```




## Linear Discriminants

## Decision Trees

```
tree.onset20=tree(intensity_level ~ newconf+onset2+onset3+onset5+onset10+onset20+duration+year_prev+dur:
summary(tree.onset20)
```

```
##
## Classification tree:
## tree(formula = intensity_level ~ newconf + onset2 + onset3 +
##       onset5 + onset10 + onset20 + duration + year_prev + duration +
##       cumulative_intensity + ep_end + incompatibility, data = war_data_valid)
## Variables actually used in tree construction:
## [1] "cumulative_intensity" "year_prev"          "newconf"
## Number of terminal nodes: 5
## Residual mean deviance: 0.5435 = 562.6 / 1035
## Misclassification error rate: 0.1288 = 134 / 1040
```

```
plot(tree.onset20)
text(tree.onset20, pretty = 0, cex=0.75)
```



## Decision Trees (With Training & Testing)

```

# Train using training set
tree.onset20=tree(intensity_level ~ newconf+onset2+onset3+onset5+onset10+onset20+duration+year_prev+dur

# Test on test set using predict()
# type="class" to return the class prediction
tree.pred=predict(tree.onset20,test_better,type="class")

# Confusion matrix
conf.matrix <- table(tree.pred,test_better$intensity_level)
conf.matrix

##
## tree.pred    0    1
##           0 142   15
##           1   15   36

# Accuracy on test set
(conf.matrix[1,1] + conf.matrix[2,2])/(conf.matrix[1,1] + conf.matrix[2,2] + conf.matrix[1,2]+ conf.mat

## [1] 0.8557692

```


## Regression Trees

```
set.seed(1)


tree.onset20=tree(intensity_level ~ newconf+onset2+onset3+onset5+onset10+onset20+duration+year_prev+dur
# Only a few of the variables were used in constructing the tree
# lstat: percentage of individuals with lower socioeconomic status
summary(tree.onset20)

##
## Classification tree:
## tree(formula = intensity_level ~ newconf + onset2 + onset3 +
##       onset5 + onset10 + onset20 + duration + year_prev + duration +
##       cumulative_intensity + ep_end + incompatibility, data = train_better)
## Variables actually used in tree construction:
## [1] "cumulative_intensity" "year_prev"          "duration"
## Number of terminal nodes:  5
## Residual mean deviance:  0.5278 = 436.5 / 827
## Misclassification error rate: 0.125 = 104 / 832


# Plot the tree
# Lower values of lstat correspond to more expensive houses
plot(tree.onset20)
text(tree.onset20,pretty=0,cex=0.75)
```



```
# cv.tree() to determine whether pruning improves performance
cv.onset20=cv.tree(tree.onset20)
# It doesn't seem to be the case
plot(cv.onset20$size,cv.onset20$dev,type="b")
```



```
# prune.tree(): function to prune to be used in case we wanted to prune the tree
prune.onset20=prune.tree(tree.onset20,best=5)
plot(prune.onset20)
text(prune.onset20,pretty=0,cex=0.75)
```



```
# Predicting based on CV results (i.e., use the unpruned tree)
yhat=predict(tree.onset20,newdata=test_better)
```

```
# plot(yhat,test_valid$onset20)
# abline(0,1)
# Test error
mse=mean((yhat-test_valid$onset20)^2)
```

```
## Warning in Ops.factor(yhat, test_valid$onset20): '-' not meaningful for factors
```

```
mse
```

```
## [1] NA
```

```
# This model leads to test predictions that are within around $5-6K of the true
# median home value for the suburb
sqrt(mse)
```

```
## [1] NA
```

## Random Forests

```

# By default randomForest() uses m=p/3 for regression and m=sqrt(p) for classification
# Let's try m=6
set.seed(1)
rf=randomForest(intensity_level ~ newconf+onset2+onset3+onset5+onset10+onset20+duration+year_prev+durat.
yhat.rf = predict(tree.onset20,newdata=test_better)

mean((yhat.rf-as.integer(test_better$intensity_level))^2)

```

```
## [1] 0.8921158
```

```

# importance(): view the importance of each variable
# %IncMSE: mean decrease of accuracy in predictions on the OOB samples when a
# given variable is excluded from the model
# IncNodeImpurity: total decrease in node impurity that results from splits over
# that variable, averaged over all trees (RSS in regr. vs. deviance in class.)
importance(rf)

```

| ##                      |           | 0          | 1         | MeanDecreaseAccuracy | MeanDecreaseGini |
|-------------------------|-----------|------------|-----------|----------------------|------------------|
| ## newconf              | 5.473129  | 13.0509594 | 10.823956 | 5.296934             |                  |
| ## onset2               | 5.752058  | 3.3409108  | 7.056920  | 1.126125             |                  |
| ## onset3               | 6.575704  | 2.5822461  | 6.976675  | 1.390664             |                  |
| ## onset5               | 8.881560  | -2.9449943 | 8.453287  | 1.952786             |                  |
| ## onset10              | 3.268522  | 5.7462269  | 6.746310  | 2.109857             |                  |
| ## onset20              | -6.250389 | 8.1246338  | 6.334236  | 2.410396             |                  |
| ## duration             | 13.288941 | 15.1205533 | 18.528238 | 15.521490            |                  |
| ## year_prev            | 24.056893 | 21.2224538 | 34.542379 | 40.501639            |                  |
| ## cumulative_intensity | 85.197794 | 82.4450055 | 91.054931 | 61.450479            |                  |
| ## ep_end               | -1.223064 | 0.7097914  | -0.587375 | 4.125935             |                  |
| ## incompatibility      | 16.503136 | 9.9642445  | 18.991830 | 10.754785            |                  |

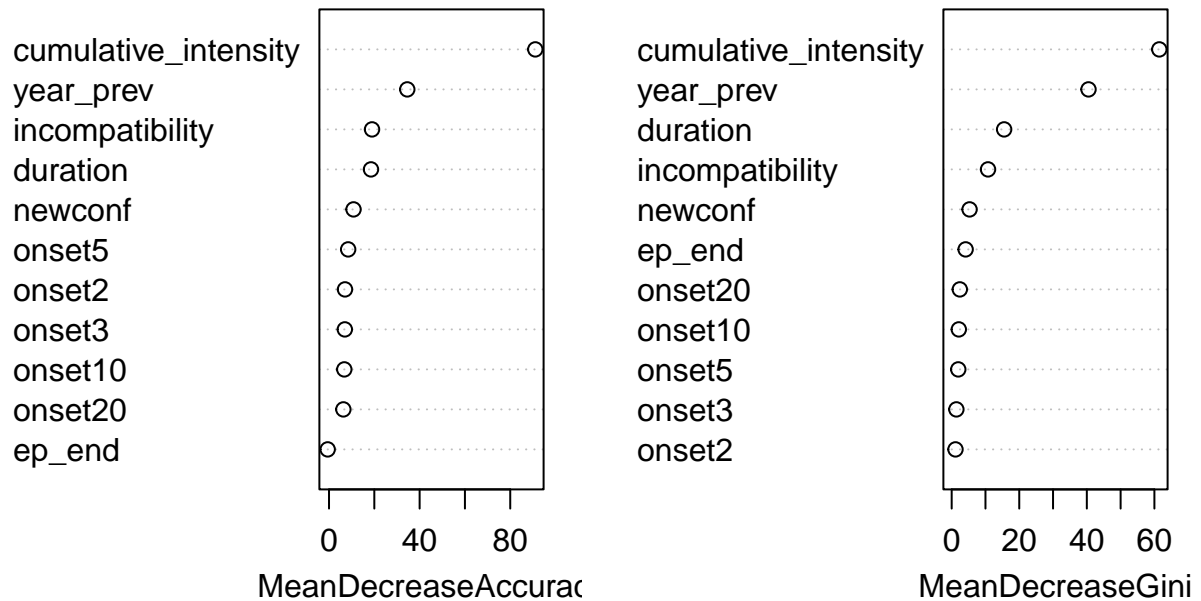
```

# varImpPlot(): Variance importance plot
varImpPlot(rf)

```



rf



## Other Models

- Penalized Logistic Regression -`plr`
- Conditional Inference Random Forest -`cforest`
- Random Forest - `rf`
- Bayesian Generalized Linear Model -`bayesglm`
- Boosted Generalized Additive Model - `gamboost`
- Support Vector Machines with Linear Kernel - `svmLinear`

```
library(caret)

#specify the cross-validation method
ctrl <- trainControl(method = "cv")

#fit a regression model and use LOOCV to evaluate performance
model <- train(intensity_level ~ newconf+onset2+onset3+onset5+onset10+onset20+duration+year_prev+duration,
               data = data,
               method = "rf",
               control = ctrl)

#view summary of LOOCV
print(model)
```

```
## k-Nearest Neighbors
##
## 832 samples
## 11 predictor
## 2 classes: '0', '1'
##
## No pre-processing
## Resampling: Cross-Validated (10 fold)
## Summary of sample sizes: 749, 749, 749, 749, 749, 749, ...
## Resampling results across tuning parameters:
##
## k Accuracy Kappa
## 5 0.8714859 0.5869409
## 7 0.8738669 0.5899490
## 9 0.8738669 0.5904556
##
## Accuracy was used to select the optimal model using the largest value.
## The final value used for the model was k = 9.
```

```
predictions <- predict(model, test_better, type="raw")
message('0 for no conflict, 1 for new conflict')
```

```
## 0 for no conflict, 1 for new conflict
```

```
message('Confusion Matrix')
```

```
## Confusion Matrix
```

```
# Confusion Matrix
table(predictions, test_better$intensity_level)
```

```
##
## predictions    0    1
##           0 142  15
##           1  15  36
```

```
# Test Error
message('Test Error Rate')
```

```
## Test Error Rate
```

```
mean(predictions != test_better$intensity_level)
```

```
## [1] 0.1442308
```

```
confusionMatrix(data = predict(model, test_better), test_better$ep_end)
```

```

## Confusion Matrix and Statistics
##
##           Reference
## Prediction  0   1
##           0 98 59
##           1 39 12
##
##           Accuracy : 0.5288
##           95% CI : (0.4586, 0.5982)
##           No Information Rate : 0.6587
##           P-Value [Acc > NIR] : 0.99996
##
##           Kappa : -0.1241
##
## Mcnemar's Test P-Value : 0.05495
##
##           Sensitivity : 0.7153
##           Specificity : 0.1690
##           Pos Pred Value : 0.6242
##           Neg Pred Value : 0.2353
##           Prevalence : 0.6587
##           Detection Rate : 0.4712
##           Detection Prevalence : 0.7548
##           Balanced Accuracy : 0.4422
##
##           'Positive' Class : 0
##

```